

Assignment #5: Deep Learning

Goal: Develop an appreciation for using **deep learning** solutions for different types of problems (regression, binary classification, multiclass classification).

Learning objectives:

- Learn how to implement deep learning solutions to classical problems in regression and (binary and multiclass) classification using Python, Keras, TensorFlow, and Jupyter Notebooks.
- Get acquainted with representative **datasets** and problems in deep learning.
- Learn how to implement, train, and evaluate **fully connected neural networks**.

Starter package

- **Starter code:** **CAP4630_A5_starter.ipynb** (available on Google Colab – check link on Canvas)

Instructions:

- **This is a group assignment.** Students are encouraged (but not required) to work in groups of max 3 students. See A4 guidelines for more details.
- Document all your findings, steps, conclusions, lessons learned, insights, etc. in your **report** (*think of it as a “lab notebook”*)
- **Add your answers to the numbered questions in (or right before) the “Conclusions” section of your report.**

Procedure:

1. Access the instructor's notebook containing the starter code on Google Colab.
NB: You must use your FAU account!
2. **Make a copy of the notebook and save it to your Google Drive.**
3. Run the starter code and ensure that it works as intended.
4. (OPTIONAL) Work on TODO items marked as “OPTIONAL” in the starter code notebook (for a **maximum of 10% bonus** on your grade for the assignment).

5. Answer the questions below:

PART 1:

- (1) What type of preprocessing was applied to the raw input data? Why was it necessary to do so?
- (2) Does your model suffer from overfitting? If so, what would you suggest doing about it?
- (3) Is accuracy a good metric of success in this case? Why (not)?

PART 2:

- (1) What type of preprocessing was applied to the raw input data? Why was it necessary to do so?
- (2) How many categories are there in this case?
- (3) Does your model suffer from overfitting? If so, what would you suggest doing about it?
- (4) Is accuracy a good metric of success in this case? Why (not)?

PART 3:

- (1) What type of preprocessing was applied to the raw input data? Why was it necessary to do so?
- (2) Why is this problem a case of *regression* (rather than *classification*)?
- (3) Does your model suffer from overfitting? If so, what would you suggest doing about it?
- (4) Is *mean absolute error* (MAE) a good metric of success in this case? Why (not)?

6. Prepare your **report**, with all relevant plots, code snippets, numerical values and – most importantly – your insights and lessons learned.
7. **Submit via Canvas the link to your Google Colab notebook** (which should automatically serve as your report – unless you choose to prepare a separate document).